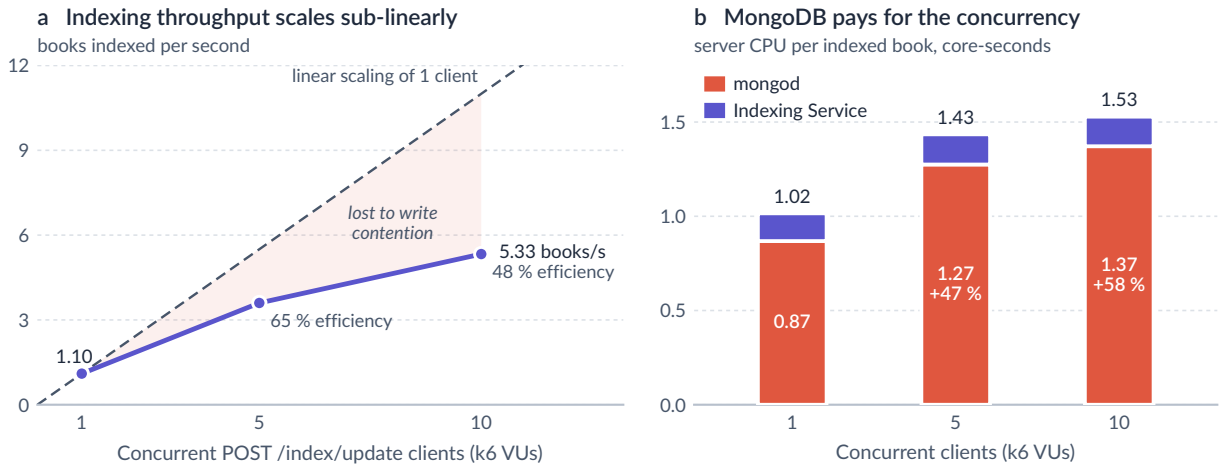
a Indexing throughput scales sub-linearly
books indexed per second
12
9
6
3
0
linear scaling of 1 client
lost to write contention
5.33 books/s
48 % efficiency
65 % efficiency
1.10
1
5
10
Concurrent POST /index/update clients (k6 VUs)
b MongoDB pays for the concurrency
server CPU per indexed book, core-seconds
mongod
Indexing Service
1.5
1.0
0.5
0.0
1.02
0.87
1.43
1.27
+47 %
1.53
1.37
+58 %
1
5
10
Concurrent clients (k6 VUs)